

Computational Linear Algebra: Floating Point Representation and Rounding Error

David Shuman

February 1, 2022

IEEE Standard Floating Point Representation

IEEE Standard Floating Point Representation

- ▶ Normalized numbers have the form

$$[\pm]1.[mantissa] * 2^{[exponent]}$$

IEEE Standard Floating Point Representation

- ▶ Normalized numbers have the form

$$[\pm]1.[mantissa] * 2^{[exponent]}$$

- ▶ Example:

$$(23.171875)_{10} = (10111.001011)_2 = +1.0111001011 * 2^4$$

IEEE Standard Floating Point Representation

- ▶ Normalized numbers have the form

$$[\pm]1.[mantissa] * 2^{[exponent]}$$

- ▶ Example:

$$(23.171875)_{10} = (10111.001011)_2 = +1.0111001011 * 2^4$$

- ▶ The leading 1 is not stored, it is implied

IEEE Standard Floating Point Representation (cont.)

- ▶ Different types of numbers have different types of bits for each part:


Precision	Sign	Exponent	Mantissa	Total
single	1	8	23	32
double	1	11	52	64
long double	1	15	64	80

IEEE Standard Floating Point Representation (cont.)

- ▶ Different types of numbers have different types of bits for each part:

Precision	Sign	Exponent	Mantissa	Total
single	1	8	23	32
double	1	11	52	64
long double	1	15	64	80

- ▶ We will almost always be using the 64-bit “double” precision:



- ▶ The sign is 0 for positive numbers and 1 for negative numbers.

IEEE Standard Floating Point Representation (cont.)

- ▶ $[\pm]1.[mantissa] * 2^{[exponent]}$
- ▶ Back to our example:

$$23.171875 = 10111.001011 = +1.0111001011 * 2^4$$

IEEE Standard Floating Point Representation (cont.)

- ▶ $[\pm]1.[mantissa] * 2^{[exponent]}$
- ▶ Back to our example:

$$23.171875 = 10111.001011 = +1.0111001011 * 2^4$$

- ▶ The sign is 0

IEEE Standard Floating Point Representation (cont.)

- ▶ $[\pm]1.[mantissa] * 2^{[exponent]}$
- ▶ Back to our example:

$$23.171875 = 10111.001011 = +1.0111001011 * 2^4$$

- ▶ The sign is 0
- ▶ The mantissa is

0111001011000

or separated into four bit chunks:

0111 0010 1100 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000

Exponents

- ▶ 11 bits allow us to represent $2^{11} = 2048$ different integers

Exponents

- ▶ 11 bits allow us to represent $2^{11} = 2048$ different integers
- ▶ We need both positive and negative exponents, and we keep two integers for special cases, so we have 2046 left to cover integers from -1022 to 1023:

binary	00000000000	00000000001	00000000010	...	11111111110	11111111111
decimal	0	1	2	...	2046	2047
exponent	special	-1022	-1021	...	1023	special

Exponents

- ▶ 11 bits allow us to represent $2^{11} = 2048$ different integers
- ▶ We need both positive and negative exponents, and we keep two integers for special cases, so we have 2046 left to cover integers from -1022 to 1023:

binary	00000000000	00000000001	00000000010	...	11111111110	11111111111
decimal	0	1	2	...	2046	2047
exponent	special	-1022	-1021	...	1023	special

- ▶ 1023 is the **exponent bias**. We add this amount to the actual exponent, and then store that number (**note: this is the same as adding 1024 and subtracting 1**)
- ▶ Subtract 1023 after storing (**or subtract 1024 and add 1**)
- ▶ Thus $+1.0111001011 * 2^4$ is

0	10000000011	011100101100...0
---	-------------	------------------

Special Cases and More Examples

Special cases:

- ▶ Exponent of 2047 is equal to ∞ if the mantissa is all zeros, and equal to NaN (Not a Number) otherwise
- ▶ Exponent of 0 used to represent subnormal (very small) numbers. See the book as we will not harp on these

Special Cases and More Examples

Special cases:

- ▶ Exponent of 2047 is equal to ∞ if the mantissa is all zeros, and equal to NaN (Not a Number) otherwise
- ▶ Exponent of 0 used to represent subnormal (very small) numbers. See the book as we will not harp on these

More examples:

1. $(-101101.0011101)_2$ is represented as

1	10000000100	01101001110100...0
---	-------------	--------------------

Special Cases and More Examples

Special cases:

- ▶ Exponent of 2047 is equal to ∞ if the mantissa is all zeros, and equal to NaN (Not a Number) otherwise
- ▶ Exponent of 0 used to represent subnormal (very small) numbers. See the book as we will not harp on these

More examples:

1. $(-101101.0011101)_2$ is represented as

1	10000000100	01101001110100...0
---	-------------	--------------------

Special Cases and More Examples

Special cases:

- ▶ Exponent of 2047 is equal to ∞ if the mantissa is all zeros, and equal to NaN (Not a Number) otherwise
- ▶ Exponent of 0 used to represent subnormal (very small) numbers. See the book as we will not harp on these

More examples:

1. $(-101101.0011101)_2$ is represented as

1	10000000100	01101001110100...0
---	-------------	--------------------

2. What do these IEEE numbers represent in binary?

1	10000110001	0111000101100...0
---	-------------	-------------------

0	01111011110	0111000101100...0
---	-------------	-------------------

Special Cases and More Examples

Special cases:

- ▶ Exponent of 2047 is equal to ∞ if the mantissa is all zeros, and equal to NaN (Not a Number) otherwise
- ▶ Exponent of 0 used to represent subnormal (very small) numbers. See the book as we will not harp on these

More examples:

1. $(-101101.0011101)_2$ is represented as

1	10000000100	01101001110100...0
---	-------------	--------------------

2. What do these IEEE numbers represent in binary?

1	10000110001	0111000101100...0
---	-------------	-------------------

0	01111011110	0111000101100...0
---	-------------	-------------------

$$-1.0111000101100 * 2^{50} \quad \text{and} \quad +1.0111000101100 * 2^{-33}$$

Rounding

Rounding

If we always **chop** (drop/ignore/truncate) the bits after b_{52} then we will systematically be rounding down. This can accumulate, so we use **rounding**:

- ▶ If what is left over is $< 1/2$:

$$1. \boxed{b_1 b_2 b_3 \cdots b_{51} b_{52}} \underbrace{0 \cdots \cdots}_{\text{round down by truncating}}$$

just drop the bits beyond position 52

- ▶ If what is left over is $> 1/2$:

$$1. \boxed{b_1 b_2 b_3 \cdots b_{51} b_{52}} \underbrace{1 \cdots \cdots}_{\text{round up by adding 1 to } b_{52}}$$

add 1 to bit b_{52} carrying as necessary

- ▶ If what is left over is $= 1/2$:

$$1. \boxed{b_1 b_2 b_3 \cdots b_{51} b_{52}} \underbrace{1 \ 0 \cdots \cdots}_{\text{all 0's}}$$

if $b_{52} = 1$: (round up)

add 1 to b_{52}

if $b_{52} = 0$: (round down)

do nothing

- ▶ in this way we always end up with $b_{52} = 0$
- ▶ these two possibilities are equally likely

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number
- ▶ In our IEEE double precision representation, this next smallest number is

$$1 + \varepsilon_M = +1.\boxed{000\dots\dots 001} = 1 + 2^{-52},$$

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number
- ▶ In our IEEE double precision representation, this next smallest number is

$$1 + \varepsilon_M = +1.\boxed{000\dots\dots 001} = 1 + 2^{-52},$$

- ▶ So $\varepsilon_M = 2^{-52}$

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number
- ▶ In our IEEE double precision representation, this next smallest number is

$$1 + \varepsilon_M = +1.\boxed{000\dots\dots\dots001} = 1 + 2^{-52},$$

- ▶ So $\varepsilon_M = 2^{-52}$
- ▶ Note: $2^{-52} = (2^{10})^{-5.2} \approx (10^3)^{-5.2} = 10^{-15.6} \approx 10^{-16}$

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number
- ▶ In our IEEE double precision representation, this next smallest number is

$$1 + \varepsilon_M = +1.\boxed{000 \dots\dots\dots 001} = 1 + 2^{-52},$$

- ▶ So $\varepsilon_M = 2^{-52}$
- ▶ Note: $2^{-52} = (2^{10})^{-5.2} \approx (10^3)^{-5.2} = 10^{-15.6} \approx 10^{-16}$
- ▶ `.Machine$double.eps`

Machine ε

- ▶ The number **machine epsilon**, denoted $\varepsilon_{\text{mach}}$ or ε_M , is the distance between 1 and the next smallest representable number
- ▶ In our IEEE double precision representation, this next smallest number is

$$1 + \varepsilon_M = +1.\boxed{000 \dots\dots\dots 001} = 1 + 2^{-52},$$

- ▶ So $\varepsilon_M = 2^{-52}$
- ▶ Note: $2^{-52} = (2^{10})^{-5.2} \approx (10^3)^{-5.2} = 10^{-15.6} \approx 10^{-16}$
- ▶ `.Machine$double.eps`
- ▶ Find gap between 2^{20} and next smallest representable number

Rounding Error

Example:



$$(9.4)_{10} = (1001.\overline{0110})_2$$

Rounding Error

Example:



$$(9.4)_{10} = (1001.\overline{0110})_2$$



$$\text{fl}(9.4) = +1. \boxed{0010\ 1100\ 1100\ 1100\ 1100\ 1100\ 1100\ 1100\ 1100\ 1100\ 1100\ 1101} * 2^3$$



So we truncated the tail by

$$.\overline{1100} * 2^{-52} * 2^3 = .\overline{0110} * 2^{-51} * 2^3 = 0.4 * 2^{-48} = 0.8 * 2^{-49}$$



And by rounding up, we added $2^{-52} * 2^3 = 1.0 * 2^{-49}$



Therefore $\text{fl}(9.4) = 9.4 + 0.2 * 2^{-49}$

Absolute Rounding Error:



$$|x_c - x| = |\text{fl}(9.4) - 9.4| = 0.2 * 2^{-49}$$

Relative Rounding Error:



$$\frac{|x_c - x|}{|x|} = \frac{|\text{fl}(9.4) - 9.4|}{9.4} = \frac{0.2 * 2^{-49}}{9.4} = \frac{8}{47} * 2^{-52}$$



The relative rounding error is always less than or equal to $\frac{\epsilon_{mach}}{2}$

Addition and Subtraction of Floating Point Numbers

- ▶ There is a register dedicated to doing computations so that the additions can be done with higher precision
- ▶ To add/subtract two numbers:
 1. Line up decimals
 2. Do the addition/subtraction (in higher precision, so you don't need to worry about rounding yet)
 3. Once the arithmetic is complete, round to convert back to floating point form