

Computational Linear Algebra: $Ax=b$ (Part III: Cholesky Factorization)

David Shuman

February 17, 2022

Review of Eigenvalues and Diagonalizability

Recall an **eigenvalue** of an $n \times n$ matrix A is a scalar $\lambda \in \mathbb{R}$ such that there exists a nonzero vector $v \in \mathbb{R}^n$ with the property that $Av = \lambda v$. The vector v is an **eigenvector** with eigenvalue λ

Review of Eigenvalues and Diagonalizability

Recall an **eigenvalue** of an $n \times n$ matrix A is a scalar $\lambda \in \mathbb{R}$ such that there exists a nonzero vector $v \in \mathbb{R}^n$ with the property that $Av = \lambda v$. The vector v is an **eigenvector** with eigenvalue λ

A is **diagonalizable** if $A = PDP^{-1}$ for some diagonal matrix D (whose diagonal elements are the eigenvalues of A) and invertible matrix P (the columns of which are eigenvectors of A)

Review of Eigenvalues and Diagonalizability

Recall an **eigenvalue** of an $n \times n$ matrix A is a scalar $\lambda \in \mathbb{R}$ such that there exists a nonzero vector $v \in \mathbb{R}^n$ with the property that $Av = \lambda v$. The vector v is an **eigenvector** with eigenvalue λ

A is **diagonalizable** if $A = PDP^{-1}$ for some diagonal matrix D (whose diagonal elements are the eigenvalues of A) and invertible matrix P (the columns of which are eigenvectors of A)

Q: When is A diagonalizable?

Review of Eigenvalues and Diagonalizability

Recall an **eigenvalue** of an $n \times n$ matrix A is a scalar $\lambda \in \mathbb{R}$ such that there exists a nonzero vector $v \in \mathbb{R}^n$ with the property that $Av = \lambda v$. The vector v is an **eigenvector** with eigenvalue λ

A is **diagonalizable** if $A = PDP^{-1}$ for some diagonal matrix D (whose diagonal elements are the eigenvalues of A) and invertible matrix P (the columns of which are eigenvectors of A)

Q: When is A diagonalizable?

A: When it has n linearly independent eigenvectors (i.e., columns of P form an **eigenbasis**)

Eigenspaces of Symmetric Matrices ($A^T = A$)

Definition

An $n \times n$ matrix A is **orthogonally diagonalizable** if there is a diagonal matrix D and an **orthogonal** matrix P (i.e., $P^T = P^{-1}$) such that

$$A = PDP^T = PDP^{-1}$$

Eigenspaces of Symmetric Matrices ($A^T = A$)

Definition

An $n \times n$ matrix A is **orthogonally diagonalizable** if there is a diagonal matrix D and an **orthogonal** matrix P (i.e., $P^T = P^{-1}$) such that

$$A = PDP^T = PDP^{-1}$$

Theorem (Spectral Theorem for Symmetric Matrices)

- ▶ *A matrix A is orthogonally diagonalizable if and only if it is symmetric*
- ▶ *A symmetric $n \times n$ matrix A has n **real** eigenvalues (counting multiplicities)*
- ▶ *The eigenspaces of a symmetric $n \times n$ matrix A are orthogonal*
- ▶ *The dimension of each eigenspace is equal to the multiplicity of its eigenvalue as a root of the characteristic equation*
- ▶ *Spectral decomposition: $A = \lambda_1 v_1 v_1^T + \lambda_2 v_2 v_2^T + \dots + \lambda_n v_n v_n^T$*
 - ▶ *Can view mapping A as a linear combination of orthogonal projections onto the different (orthogonal) eigenspaces, with the weights determined by the eigenvalues*

Symmetric Positive Definite Matrices

Note: If A is an $n \times n$ matrix and $x \in \mathbb{R}^n$, then $x^T A x$ is a scalar

Definition

An $n \times n$ real matrix A is **positive definite** if $x^T A x > 0$ for all vectors $x \in \mathbb{R}^n$ with $x \neq 0$ (and **positive semidefinite** if $x^T A x \geq 0$)

Theorem

- ▶ *A symmetric matrix A is positive definite if and only if all of its eigenvalues are (strictly) positive*
- ▶ *Any principal submatrix of a symmetric positive-definite matrix is symmetric positive-definite*

Symmetric positive definite matrices show up *a lot* in applications

Symmetric Positive Definite Matrices

Example

- If $A_1 = \begin{pmatrix} 3 & 2 \\ 2 & 6 \end{pmatrix}$ then

$$(x_1 \ x_2) \begin{pmatrix} 3 & 2 \\ 2 & 6 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = 3x_1^2 + 4x_1x_2 + 6x_2^2 = x_1^2 + 2(x_1 + x_2)^2 + 4x_2^2$$

so A_1 is positive definite

- Find the eigenvalues:

$$\begin{pmatrix} 8 & 4 & 2 \\ 4 & 6 & 0 \\ 2 & 0 & 3 \end{pmatrix}, \begin{pmatrix} 3 & 2 & 0 \\ 2 & 3 & 3 \\ 0 & 2 & 3 \end{pmatrix}, \begin{pmatrix} 3 & -1 & 0 & 0 & 0 & .5 \\ -1 & 3 & -1 & 0 & .5 & 0 \\ 0 & -1 & 3 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & 0 \\ 0 & .5 & 0 & -1 & 3 & -1 \\ .5 & 0 & 0 & 0 & -1 & 3 \end{pmatrix}$$

- Symmetric but not positive definite:

$$(1 \ -1) \begin{pmatrix} 1 & 2 \\ 2 & 1 \end{pmatrix} \begin{pmatrix} 1 \\ -1 \end{pmatrix} = -2 < 0$$

Cholesky Factorization

- ▶ Assume that A is symmetric positive definite
- ▶ LU factorization does not take full advantage of the structure of A
- ▶ LU factorization also does not always return L and U matrices that are symmetric (or even symmetric on the off-diagonal elements)
 - ▶ Example:

$$A = \begin{pmatrix} 4 & -2 & 0 \\ -2 & 2 & -3 \\ 0 & -3 & 10 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 4 & -2 & 0 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = LU$$

Cholesky Factorization

- ▶ Assume that A is symmetric positive definite
- ▶ LU factorization does not take full advantage of the structure of A
- ▶ LU factorization also does not always return L and U matrices that are symmetric (or even symmetric on the off-diagonal elements)

▶ Example:

$$A = \begin{pmatrix} 4 & -2 & 0 \\ -2 & 2 & -3 \\ 0 & -3 & 10 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 4 & -2 & 0 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = LU$$

- ▶ Enter the Cholesky factorization: $A = R^T R$, where R is upper-triangular

Cholesky Factorization

- ▶ Assume that A is symmetric positive definite
- ▶ LU factorization does not take full advantage of the structure of A
- ▶ LU factorization also does not always return L and U matrices that are symmetric (or even symmetric on the off-diagonal elements)

▶ Example:

$$A = \begin{pmatrix} 4 & -2 & 0 \\ -2 & 2 & -3 \\ 0 & -3 & 10 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 4 & -2 & 0 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = LU$$

- ▶ Enter the Cholesky factorization: $A = R^T R$, where R is upper-triangular
- ▶ Leverages the structure of A to require half as many operations as LU

Cholesky Factorization

- ▶ Assume that A is symmetric positive definite
- ▶ LU factorization does not take full advantage of the structure of A
- ▶ LU factorization also does not always return L and U matrices that are symmetric (or even symmetric on the off-diagonal elements)

▶ Example:

$$A = \begin{pmatrix} 4 & -2 & 0 \\ -2 & 2 & -3 \\ 0 & -3 & 10 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 4 & -2 & 0 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = LU$$

- ▶ Enter the Cholesky factorization: $A = R^T R$, where R is upper-triangular
- ▶ Leverages the structure of A to require half as many operations as LU
- ▶ Every symmetric positive definite A has a unique $R^T R$ factorization

Cholesky Factorization

- ▶ Assume that A is symmetric positive definite
- ▶ LU factorization does not take full advantage of the structure of A
- ▶ LU factorization also does not always return L and U matrices that are symmetric (or even symmetric on the off-diagonal elements)

▶ Example:

$$A = \begin{pmatrix} 4 & -2 & 0 \\ -2 & 2 & -3 \\ 0 & -3 & 10 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -0.5 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 4 & -2 & 0 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = LU$$

- ▶ Enter the Cholesky factorization: $A = R^T R$, where R is upper-triangular
- ▶ Leverages the structure of A to require half as many operations as LU
- ▶ Every symmetric positive definite A has a unique $R^T R$ factorization
- ▶ Cholesky factorization provides a way to tell if a symmetric matrix is positive definite
- ▶ Example and general method on board

Cholesky Factorization: Pseudocode and Complexity

Pseudocode:

$$R = A$$

for $k = 1$ to n

for $j = k + 1$ to n

$$R_{j,j:n} = R_{j,j:n} - R_{k,j:n} * \frac{R_{k,j}}{R_{k,k}}$$

$$R_{k,k:n} = \frac{1}{\sqrt{R_{k,k}}} * R_{k,k:n}$$

Cholesky Factorization: Pseudocode and Complexity

Pseudocode:

$$R = A$$

for $k = 1$ to n

for $j = k + 1$ to n

$$R_{j,j:n} = R_{j,j:n} - R_{k,j:n} * \frac{R_{k,j}}{R_{k,k}}$$

$$R_{k,k:n} = \frac{1}{\sqrt{R_{k,k}}} * R_{k,k:n}$$

Complexity:

$O(\frac{n^3}{3})$ additions and multiplications (half as many as LU)

Cholesky Factorization: Stability and Solving $Ax = b$

► Stability

- Unlike Gaussian elimination, this algorithm is always stable in the sense that $R^T R$ is close to A
- Intuitive reason: Norms of factors cannot grow too large (recall the examples we showed where L or U had very large matrix norms in comparison to the norm of A)
- Example in 2-norm: $\|R\|_2 = \|R^T\|_2 = \sqrt{\|A\|_2}$
- No need for pivoting thanks to positive definite structure

Cholesky Factorization: Stability and Solving $Ax = b$

► Stability

- Unlike Gaussian elimination, this algorithm is always stable in the sense that $R^T R$ is close to A
- Intuitive reason: Norms of factors cannot grow too large (recall the examples we showed where L or U had very large matrix norms in comparison to the norm of A)
- Example in 2-norm: $\|R\|_2 = \|R^T\|_2 = \sqrt{\|A\|_2}$
- No need for pivoting thanks to positive definite structure



► Solve $Ax = R^T R x = b$ in two steps:

- $R^T y = b$
- $R x = y$
- Each of these two triangular systems is $O(n^2)$ [same as LU]


Sparse Matrices

- ▶ In *many* applications, the matrices we need are **sparse**: a large number of the entries are 0
 - ▶ Non-sparse matrices are called **full** or **dense** matrices
- ▶ Typically a sparse $n \times n$ matrix A has $\leq O(n)$ nonzero entries
- ▶ Here are a few examples (nonzero entries shown in black)

$$\begin{bmatrix} 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 2 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 3 & 3 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 4 & 5 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 5 & 8 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 6 & 13 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 7 & 21 \end{bmatrix},$$




- ▶ Sometimes the nonzero entries have patterns, like in banded matrices:



$$\begin{pmatrix} A & I & O & \cdots & O & O \\ I & A & I & \cdots & O & O \\ O & I & A & \cdots & O & O \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ O & O & O & \cdots & I & A \end{pmatrix}$$

Sparse Matrices, cont.



- ▶ The idea is to record the list of positions of nonzero entries (i_p, j_p) and the matrix entry a_p in that position (a **triplet** of information)

$$\begin{array}{l} \text{positions} \begin{bmatrix} i_1 & i_2 & i_3 & i_4 & \cdots & i_n \\ j_1 & j_2 & j_3 & j_4 & \cdots & j_n \end{bmatrix} \\ \text{entries} \begin{bmatrix} a_1 & a_2 & a_3 & a_4 & \cdots & a_n \end{bmatrix} \end{array}$$

- ▶ i.e., do not store the 0's
- ▶ Then you need to get your matrix algorithms to work correctly with these new data structures
- ▶ Can be even more efficient if there is a strong pattern in the nonzero entries as in banded matrices
 - ▶ Just store the bands and the positions are understood

Sparse Matrices in R

- ▶ Need to include the **Matrix** package

- ▶ Try

```
A = Matrix(0,nrow=10,ncol=10,sparse=TRUE)
A[1,3] = A[5,6] = A[10,1] = 1
A[10,10] = A[2,10] = A[5,10] = A[7,2] = 2
A
image(A)
```

- ▶ In triplet form:

```
A = spMatrix(10,20, i = c(1,3:8), j = c(2,9,6:10), x = 7 * (1:7))
A
image(A)
```

- ▶ Try computing

```
3 * A
A %*% rep(0,20)
t(A)
A %*% t(A)
A + 1
```

Sparse Matrices in R

► Saving Space

```
m1 = matrix(0, nrow = 1000, ncol = 1000)
m2 = Matrix(0, nrow = 1000, ncol = 1000, sparse = TRUE)
m3 = spMatrix(nrow = 1000, ncol = 1000)
object.size(m1)
object.size(m2)
object.size(m3)
```

► m1[500, 500] = 1 m2[500, 500] = 1 m3[500, 500] = 1 object.size(m1) object.size(m2) object.size(m3)

Sparse Matrices in R

Saving Time

- ▶ Matrix multiplication for sparse matrices is optimized to take advantage of the fact that the matrix is sparse
- ▶ For example, the matrix-vector product Ax takes $O(n^2)$ multiplications if the matrix A is $n \times n$ and dense
- ▶ If A has only $O(n)$ nonzero entries then Ax can be optimized to take only $O(n)$ multiplications